

# **CUDA**

## **Game of Life parallelization**

**Failla Edoardo**


# Introduction

Conway's game of life game → 2D grid of cellular automata

Cell state  $\begin{cases} 1 & \text{- if cell is alive} \\ 0 & \text{- if cell is dead} \end{cases}$


At each generation we apply cell transition rules:

## Birth




## Death

- Overcrowding:



- Exposure:



## Survival






# Introduction

## Complexity of Conway's Game of Life

- **Time complexity:**  $O(G \times N \times M)$ , where  $G$  is the number of generations,  $M$  and  $N$  the size of the grid
- **Memory Complexity:**  $O(N \times M)$ , where  $N \times M$  is the grid size.

### Challenges:

- **Incomplete Neighbors:**
  - **Corners:** only 3 neighbors instead of 8.
  - **Edges:** only 5 neighbors.



# CUDA implementation

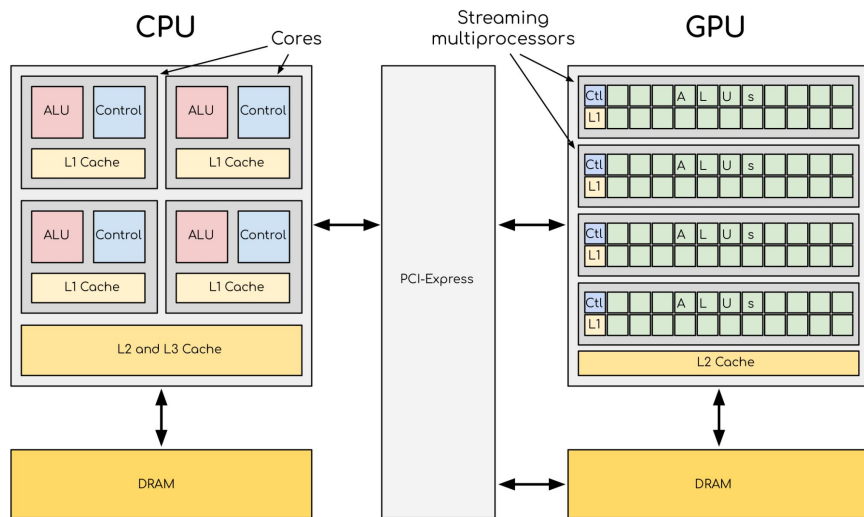
## Using GPGPU computation

### We can benefit from:

- *Natural parallelization architecture (SIMD)*
- *High Throughput*
- *Hardware Acceleration*
- *Scalability*

### CUDA (Compute Unified Device Architecture):

- A parallel computing platform
- Developed by NVIDIA for GPU processing
- Supports languages like C, C++, and Fortran.



# CUDA implementation

## Thread indexing

Game grid 2D  $\longrightarrow$  1D  $\longrightarrow$  **Cell\_index** = row x N + column


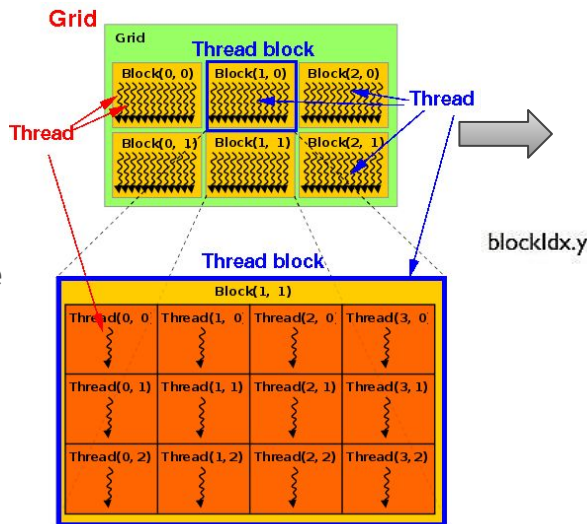
Global indexes with **Block/Grid** built-in variable

Device Tesla K40m:

- Max threads per block: 1024
- Max dimension (1024, 1024, 64)

We tested different dimensions of Block size

- (4, 4, 1)
- (8, 8, 1)
- (16, 16, 1)
- (32, 32, 1)



# CUDA implementation

## Input and parameters

```
./game_of_life <initial_state> <grid-size> <BlockDim> <Num-of-generations> --options
```

list options

--verbose "print result in a .txt file"

--check "(for one version only) check the correctness of the result"

game\_of\_life.cu

```
...  
dim3 blockDim(blockDimX, blockDimY);  
dim3 gridDim((N + blockDim.x - 1) / blockDim.x, (N + blockDim.y - 1) / blockDim.y);
```

*CUDA grid dimensions calculated*

# CUDA implementation

## First idea:


- Only one kernel
- **for** generation cycle inside kernel
- One function call



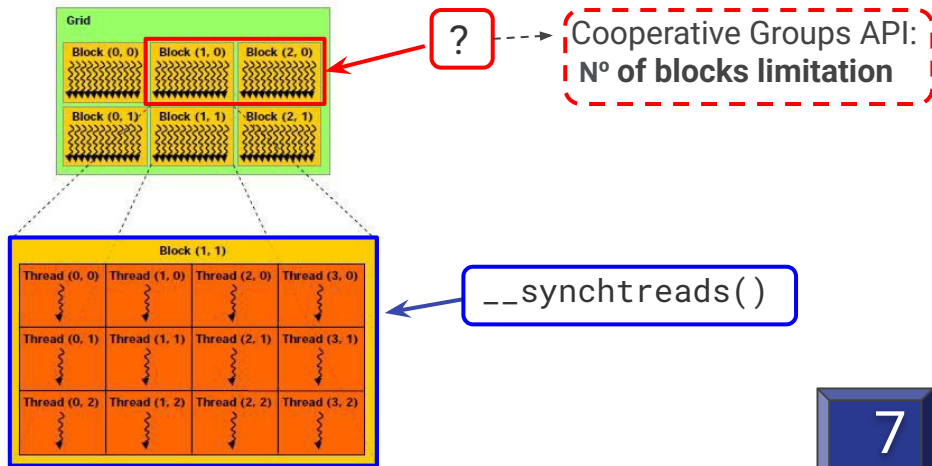
## Problems:

- Threads synchronization
- Race condition

## Race condition

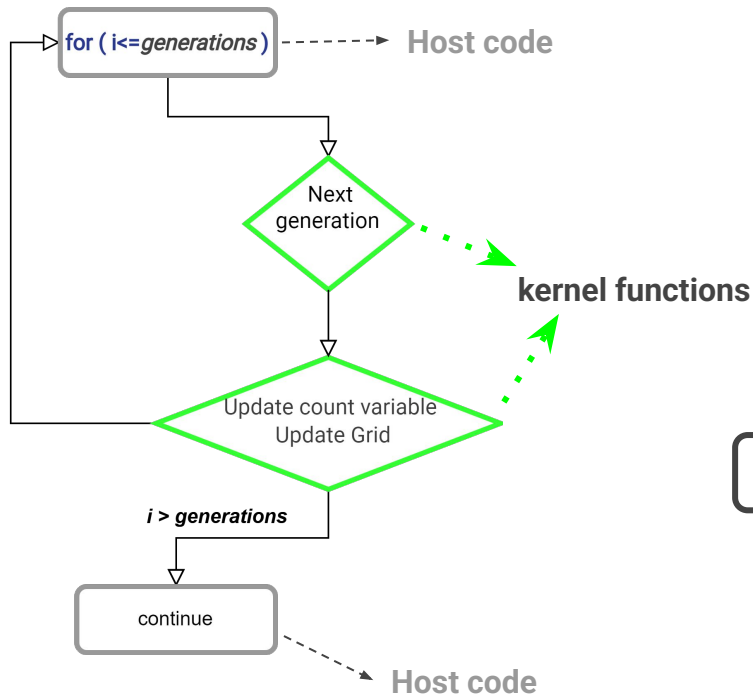


## Synchronization has many limitations




# CUDA implementation

## First implementation



## Calculate next generation:

- Check if neighbors' indexes are in range
- Count alive neighbors
- Apply rules






# CUDA implementation

## Second implementation

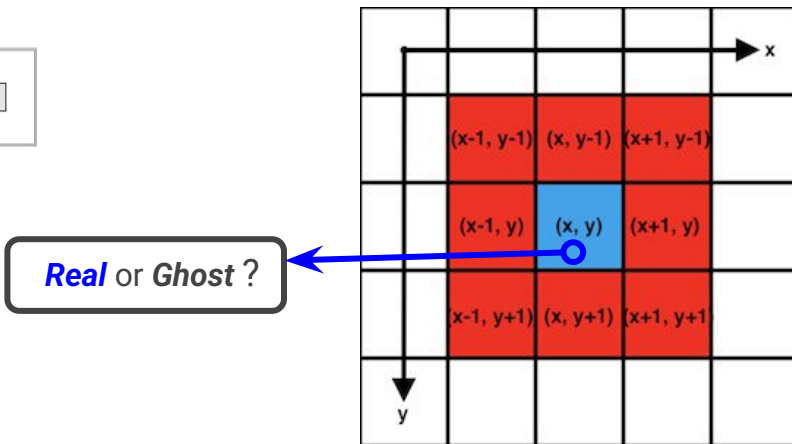
### Add Ghost cells to perimeter:

- **Ghost** cells are always **dead**
- Do not influence the rules



### Calculate next generation:

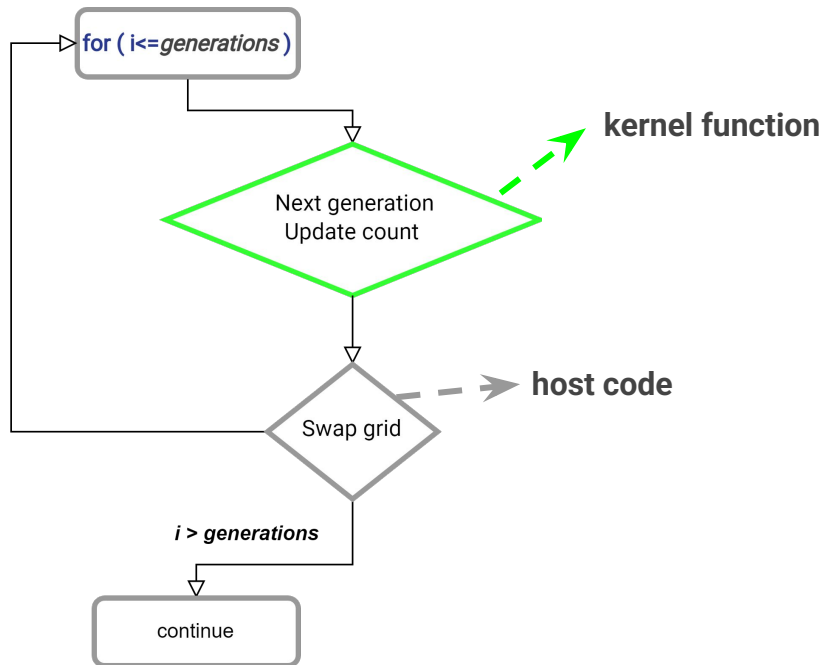
- Check if cells are **Real** or **Ghost**
- Count alive neighbors
- Apply rules



# CUDA implementation

## Some improvements

We can modify the code for optimizing grid swapping



- Only one kernel
- Reduced function call
- Swap pointer operation

# CUDA implementation

## Can we use multiple GPUs ?



### HACTAR

HACTAR è un cluster InfiniBand da oltre 20 TFLOPS con le seguenti caratteristiche:

Architecture	Linux Infiniband-QDR MIMD Distributed Shared-Memory
Node Interconnect	Infiniband QDR 40 Gb/s
Service Network	Gigabit Ethernet 1 Gb/s
CPU Model	2x Intel Xeon E5-2680 v3 2.50 GHz 12 cores
GPU Model	2x nVidia Tesla K40 - 12 GB - 2880 cuda cores
Sustained performance (Rmax)	20.13 TFLOPS (last update: june 2018)
Peak performance (Rpeak)	25.61 TFLOPS (last update: june 2018)



We can use 2 GPUs for one node

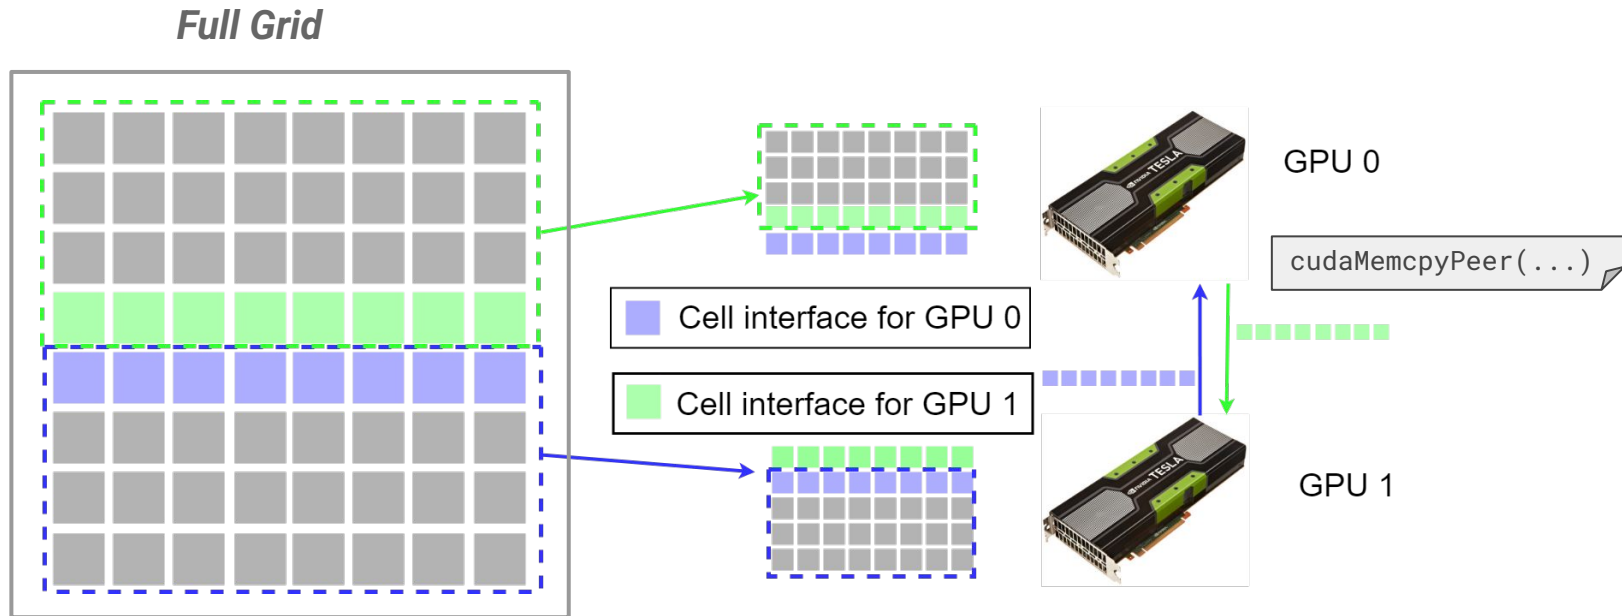
### We need to:

1. Split the grid
2. Exchange border information
3. Reassemble the grid

How can we do such task?

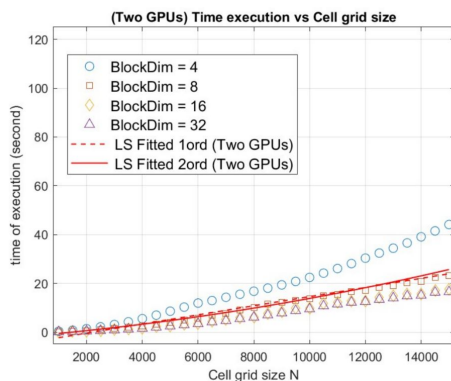
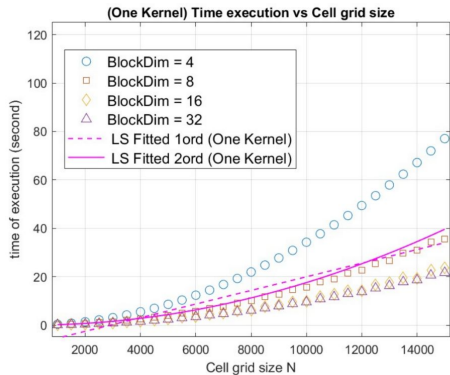
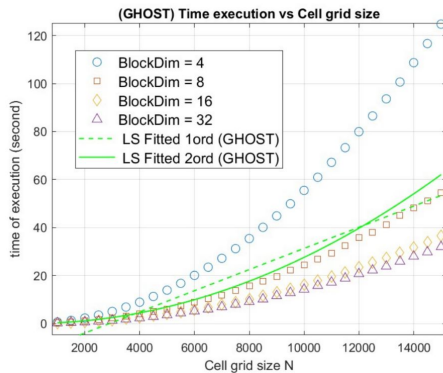
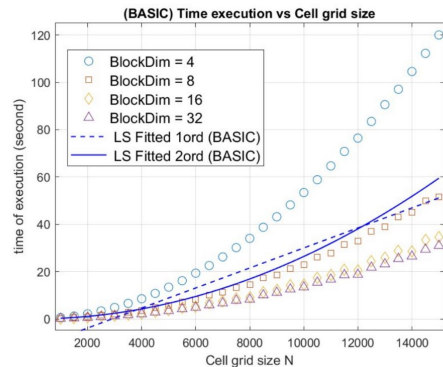
# CUDA implementation

## Two-GPUs

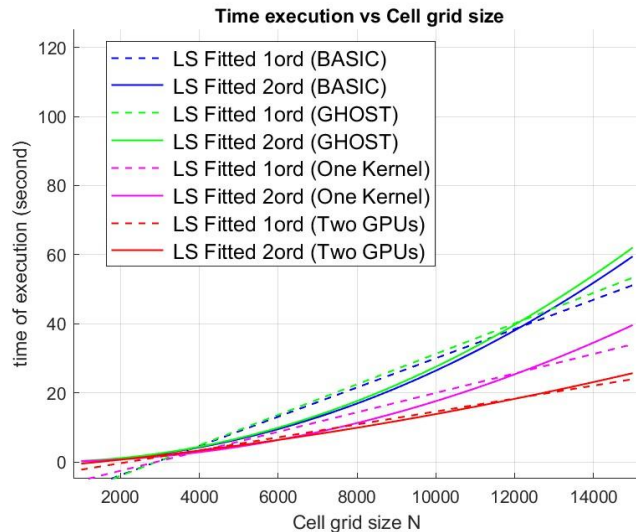


# CUDA implementation

## Time of execution



## Results Summary



# CUDA implementation

## CUDA occupancy

**“Occupancy”** is defined as the ratio of active warps on a SM to the maximum number of active warps supported by the SM.”




“The SM has a maximum number of **warps** that can be active at once.”

“The SM has a maximum number of **blocks** that can be active at once.”



“The multiprocessor creates, manages, schedules, and executes threads in groups of 32 parallel threads called **warps**”



Kernel: kernelHitBlockLimit		Grid Dim: {240, 1, 1} 240		Block Dim: {1, 1, 1} 1	
Device: Quadro K6000		Compute Capability: 3.5 Dyn Shm/Block: 0		Stat Shm/Block: 0	
Variable	Achieved	Theoretical	Device Limit		
Occupancy Per SM					
Active Blocks		16	16		
Active Warps	15.93	16	64		
Active Threads		16	2048		
Occupancy	24.90 %	25.00 %	100.00 %		

(NVIDIA doc. example)

# CUDA implementation

## CUDA occupancy

Using the GPU on my laptop

```
ncu --set full ./0kern_game_of_life <...>
```

### Section: Occupancy

```
-----  
Block Limit SM                block  
Block Limit Registers         block  
Block Limit Shared Mem       block  
Block Limit Warps             block  
Theoretical Active Warps per SM  warp  
→ Theoretical Occupancy        %  
→ Achieved Occupancy           %  
Achieved Active Warps Per SM    warp  
-----
```

BlockDim(4,4)

```
-----  
16  
64  
100  
48  
16  
● 33,33  
● 30,55  
14,66  
-----
```

BlockDim(8,8)

```
-----  
16  
32  
100  
24  
32  
● 66,67  
● 59,10  
28,37  
-----
```

BlockDim(16,16)

```
-----  
16  
8  
100  
6  
48  
● 100  
● 83,29  
39,98  
-----
```

BlockDim(32,32)

```
-----  
16  
2  
100  
1  
32  
● 66,67  
● 53,85  
25,85  
-----
```



***Thank you for the  
attention!***